

Conformal Correlation Counterexamples

Peter Cotton

`peter.cotton@microprediction.com`

September 2026

Abstract

The conformal correlation matrix of Perlo et al. (2026) is the phi coefficient between indicators of class membership in conformal prediction sets, offered as a map of which classes a classifier confuses. Seven hand-checkable examples show that it is not: a perfect classifier and a perfectly wrong one share a matrix, a zero entry coexists with 99% directional confusion, and pooling two clients with zero correlation on each produces an entry of ± 0.64 . The reason is that the matrix depends only on the law of the prediction set, while confusion depends on the joint law of the set, the true label and the ranked prediction. Two lemmas give every entry for classes that never share a set as a function of the inclusion rates alone, and recover every inclusion rate exactly from three entries when sets are singletons, which bears on the privacy claim made for the matrix.

1 The statistic

There are H classes. A classifier maps an input X to a probability vector $p(X)$ and to a prediction set $C(X) \subseteq \{1, \dots, H\}$ built by split conformal prediction (Vovk et al., 2005; Lei et al., 2018), in the paper by the adaptive prediction sets of Romano et al. (2020). The true label is Y and the class the classifier ranks first is $\hat{Y}$. For each class h define the membership indicator

$$z_h = \mathbf{1}\{h \in C(X)\},$$

and write $p_h = \mathbb{E}[z_h]$ for the inclusion rate of class h and $q_{ij} = \mathbb{E}[z_i z_j]$ for the co-occurrence rate of classes i and j . The conformal correlation matrix (CCM) of Perlo et al. (2026) has entries

$$\rho_{ij} = \frac{\mathbb{E}[z_i z_j] - \mathbb{E}[z_i] \mathbb{E}[z_j]}{\sqrt{\text{Var}(z_i) \text{Var}(z_j)}} = \frac{q_{ij} - p_i p_j}{\sqrt{p_i(1-p_i) p_j(1-p_j)}}, \quad (1)$$

which is the Pearson correlation of two Bernoulli variables, also called the phi coefficient (Yule, 1912). In practice the expectations are sample averages over a test set.

The authors attach three readings to the entries. Their Remark 1 says that intrinsically similar classes “might be expected to exhibit high positive correlation”. Their Remark 2 says that pairs of dissimilar classes “are expected to display strong negative correlation”. Their Remark 3 says that classes “unrelated or independent with respect to the classifier representation are expected to yield correlation coefficients close to zero”. On the strength of these readings they conclude that a positive entry between two classes means the classifier “is much more likely to misclassify” one as the other, that the matrix “contains a superset of the information conveyed by the confusion matrix”, and that it “captures the likelihood of class pairs co-occurring within conformal prediction sets”. Each of these statements is false, and the examples below show it.

2 What the matrix can and cannot see

Everything below follows from one observation.

Proposition 1. *The conformal correlation matrix is a functional of the law of the prediction set $C(X)$ alone. Two situations with the same law of $C(X)$ have the same matrix, whatever the joint law of $(C(X), Y, \hat{Y})$.*

Proof. Every quantity in (1) is an expectation of a function of the vector $(z_1, \dots, z_H)$, and that vector is a function of $C(X)$. $\square$

Confusion is a different object. The confusion matrix has entries $\mathbb{P}(Y = a, \hat{Y} = b)$, which depend on the joint law of the true label and the ranked prediction. The set $C(X)$ can be held fixed while Y or $\hat{Y}$ is changed freely, subject only to whatever coverage constraint is imposed, and by Proposition 1 the matrix cannot notice.

The coverage constraint is weak. Split conformal prediction guarantees, under exchangeability of calibration and test data,

$$\mathbb{P}(z_Y = 1) = \sum_y \mathbb{P}(Y = y, z_y = 1) \geq 1 - \alpha, \quad (2)$$

which is one scalar. It says nothing about q_{ij} , which mixes the true class with a wrong class, two wrong classes, and different true classes on different inputs.

Proposition 2. *For every $\rho \in [-1, 1]$ there is a classifier with coverage one and accuracy one whose matrix has an off-diagonal entry equal to ρ .*

Proof. Let the true label always be T , always in the set and always ranked first, so coverage and accuracy are one. Let two other labels A and B be added to the set according to

$$\begin{aligned} \mathbb{P}(z_A = 0, z_B = 0) &= \mathbb{P}(z_A = 1, z_B = 1) = \frac{1 + \rho}{4}, \\ \mathbb{P}(z_A = 1, z_B = 0) &= \mathbb{P}(z_A = 0, z_B = 1) = \frac{1 - \rho}{4}. \end{aligned}$$

Then $p_A = p_B = \frac{1+\rho}{4} + \frac{1-\rho}{4} = \frac{1}{2}$, so each variance is $\frac{1}{4}$, and $q_{AB} = \frac{1+\rho}{4}$. Substituting into (1),

$$\rho_{AB} = \frac{(1 + \rho)/4 - 1/4}{1/4} = \rho. \quad \square$$

So the whole correlation scale is available with no classification problem between A and B at all. The examples that follow are less abstract than this, and each contradicts a specific sentence in the paper.

3 Seven examples

Each example needs only the inclusion rates p_i and the co-occurrence rates q_{ij} , substituted into (1). We order them from the simplest to refute and the most absurd result downward. In every example coverage is one, so the certificate is satisfied with room to spare.

Example 1 (A perfect classifier and a perfectly wrong one share a matrix). Three classes A, B, C with equal frequency, and sets that always contain the true class and one neighbor.

True class	Prediction set	Good model ranks first	Bad model ranks first
A	$\{A, B\}$	A	B
B	$\{B, C\}$	B	C
C	$\{C, A\}$	C	A

Each label is in two of the three sets, so every $p_i = 2/3$ and every variance is $2/9$. Each pair shares exactly one of the three sets, so every $q_{ij} = 1/3$. Every off-diagonal entry is

$$\rho_{ij} = \frac{1/3 - (2/3)(2/3)}{2/9} = -\frac{1}{2}.$$

The good model has accuracy one. The bad model has accuracy zero. The labels, the sets and therefore the matrix are identical for both. The confusion matrix of the good model is the identity and that of the bad model has an empty diagonal, so the matrix does not contain “a superset of the information conveyed by the confusion matrix”.

Example 2 (Zero correlation with 99% directional confusion). Three classes and a small number ε , say 0.01. Inputs arrive in four kinds.

Probability	True class	Prediction set	Ranked first
$(1 - \varepsilon)/2$	A	$\{A, B\}$	B
$\varepsilon/2$	A	$\{A\}$	A
$(1 - \varepsilon)/2$	B	$\{B\}$	B
$\varepsilon/2$	C	$\{C\}$	C

A is in the sets of the first two kinds, so $p_A = (1 - \varepsilon)/2 + \varepsilon/2 = 1/2$. B is in the sets of the first and third kinds, so $p_B = 1 - \varepsilon$. They appear together only in the first kind, so $q_{AB} = (1 - \varepsilon)/2 = p_{APB}$. The numerator of (1) is zero, so

$$\rho_{AB} = 0.$$

Among inputs whose true class is A , the classifier ranks B first with probability $1 - \varepsilon$, which is 99% at $\varepsilon = 0.01$. The authors’ Remark 3 reads a zero entry as classes “unrelated or independent with respect to the classifier representation”.

Example 3 (Perfect correlation with one in a million co-occurrence). Every set contains the true class T , ranked first. Two other labels A and B are added to the set together with probability ε , and otherwise neither is added. Then $z_A = z_B$ on every input, so

$$\rho_{AB} = 1 \quad \text{for every } 0 < \varepsilon < 1,$$

while the probability that A and B appear together in a set is $q_{AB} = \varepsilon$, which may be 0.4, or 0.01, or one in a million. The abstract of the paper says the matrix “captures the likelihood of class pairs co-occurring”. The likelihood is ε . The entry is one regardless.

Example 4 (Every pair positive from difficulty alone). Three classes with equal frequency. Half the inputs are hard and receive the set of all three classes. The other half are easy and receive the singleton set of their true class. The true class is ranked first in both cases, so accuracy is one. Each label is in every hard set and in a third of the easy sets, so

$$p_i = \frac{1}{2} \cdot 1 + \frac{1}{2} \cdot \frac{1}{3} = \frac{2}{3}, \quad p_i(1 - p_i) = \frac{2}{9}.$$

Two labels appear together only in hard sets, so $q_{ij} = 1/2$ for every pair, and

$$\rho_{ij} = \frac{1/2 - (2/3)(2/3)}{2/9} = \frac{1}{4}.$$

Every pair is positively correlated and the classifier is perfect. The correlation has one cause, whether the input was hard, and says nothing about any pair. The authors’ Remark 1 reads a positive entry as “intrinsically similar classes”.

Example 5 (Pooling two clients creates a strong correlation from nothing). Every set contains the true class T , ranked first. On client 1 the labels A and B are each added to the set independently with probability 0.9. On client 2 they are each added independently with probability 0.1. The clients contribute equally to the pooled data. Within each client z_A and z_B are independent, so $\rho_{AB} = 0$ on each. Pooled, $p_A = p_B = (0.9 + 0.1)/2 = 1/2$, each variance is $1/4$, and

$$q_{AB} = \frac{1}{2}(0.9 \times 0.9) + \frac{1}{2}(0.1 \times 0.1) = 0.41, \quad \rho_{AB} = \frac{0.41 - 1/4}{1/4} = 0.64.$$

Reverse the rates on one client, so that client 1 adds A with probability 0.9 and B with probability 0.1 and client 2 does the opposite. Then $q_{AB} = \frac{1}{2}(0.09) + \frac{1}{2}(0.09) = 0.09$ and $\rho_{AB} = (0.09 - 0.25)/0.25 = -0.64$. The paper computes matrices across clients that are non identically distributed by construction and reads the entries as properties of the classifier. A strong entry of either sign is produced here by the mixture of clients alone.

Example 6 (Two classes the model never confuses get different entries). Three classes, cat, dog and plane, with equal frequency. The classifier returns the same set for every input of a given class.

True class	Prediction set	z_{cat}	z_{dog}	z_{plane}
cat	{cat, dog}	1	1	0
dog	{dog}	0	1	0
plane	{plane}	0	0	1

A third of the inputs are cats, so $p_{\text{cat}} = 1/3$, and likewise $p_{\text{dog}} = 2/3$ and $p_{\text{plane}} = 1/3$. Each variance is $2/9$. Cat and dog appear together on the cat inputs, so $q_{\text{cat,dog}} = 1/3$, and plane never appears with either, so the other two co-occurrence rates are zero. Substituting,

$$\begin{aligned} \rho_{\text{cat,dog}} &= \frac{1/3 - (1/3)(2/3)}{2/9} = \frac{1}{2}, \\ \rho_{\text{cat,plane}} &= \frac{0 - (1/3)(1/3)}{2/9} = -\frac{1}{2}, \\ \rho_{\text{dog,plane}} &= \frac{0 - (2/3)(1/3)}{2/9} = -1. \end{aligned}$$

The authors’ Remark 2 reads this as dog and plane being strongly dissimilar and cat and plane half as dissimilar. But the classifier treats plane identically with respect to both. It never puts plane in a set with either. The two entries differ only because dog is in twice as many sets as cat. Lemma 1 gives the general statement.

Example 7 (One relocated observation moves an entry from one to almost zero). There are n test inputs, and two rare labels A and B each appear in exactly one set. If they appear in the same set then $z_A = z_B$ and $\rho_{AB} = 1$. If B appears in a different set then $p_A = p_B = 1/n$, $q_{AB} = 0$, and

$$\rho_{AB} = \frac{0 - 1/n^2}{(1/n)(1 - 1/n)} = -\frac{1}{n-1},$$

which at $n = 1000$ is about -0.001 . The paper presents heat maps of entries with no uncertainty attached. For rare labels a single observation moves an entry across its whole range.

4 Two lemmas

Write $o_i = p_i/(1 - p_i)$ for the odds of inclusion of class i .

Lemma 1 (The floor, and the attainable range). *If classes i and j never share a set, then*

$$\rho_{ij} = -\sqrt{o_i o_j} = -\sqrt{\frac{p_i p_j}{(1 - p_i)(1 - p_j)}}. \quad (3)$$

With H classes of equal inclusion rate and singleton sets this equals $-1/(H - 1)$. More generally, for fixed inclusion rates $p_i \leq p_j$ the entry ρ_{ij} ranges over

$$\left[-\sqrt{o_i o_j}, \sqrt{o_i o_j} \right] \quad \text{if } p_i + p_j \leq 1, \quad \left[-1/\sqrt{o_i o_j}, \sqrt{o_i o_j} \right] \quad \text{if } p_i + p_j > 1.$$

Proof. If $q_{ij} = 0$ then (1) reads

$$\rho_{ij} = \frac{-p_i p_j}{\sqrt{p_i(1 - p_i)p_j(1 - p_j)}} = -\sqrt{\frac{p_i^2 p_j^2}{p_i(1 - p_i)p_j(1 - p_j)}} = -\sqrt{\frac{p_i p_j}{(1 - p_i)(1 - p_j)}}.$$

With singleton sets every pair has $q_{ij} = 0$, and with equal inclusion rates $p_i = 1/H$, so $o_i = 1/(H - 1)$ and $\rho_{ij} = -1/(H - 1)$.

For the range, q_{ij} is the probability of the intersection of two events with probabilities $p_i \leq p_j$, so by the Fréchet bounds $\max(0, p_i + p_j - 1) \leq q_{ij} \leq p_i$. At the upper bound the numerator of (1) is $p_i - p_i p_j = p_i(1 - p_j)$ and

$$\rho_{ij} = \frac{p_i(1 - p_j)}{\sqrt{p_i(1 - p_i)p_j(1 - p_j)}} = \sqrt{\frac{p_i(1 - p_j)}{(1 - p_i)p_j}} = \sqrt{o_i o_j}.$$

At the lower bound, if $p_i + p_j \leq 1$ then $q_{ij} = 0$ and (3) applies. If $p_i + p_j > 1$ then $q_{ij} = p_i + p_j - 1$, the numerator is $p_i + p_j - 1 - p_i p_j = -(1 - p_i)(1 - p_j)$, and

$$\rho_{ij} = \frac{-(1 - p_i)(1 - p_j)}{\sqrt{p_i(1 - p_i)p_j(1 - p_j)}} = -\sqrt{\frac{(1 - p_i)(1 - p_j)}{p_i p_j}} = -1/\sqrt{o_i o_j}. \quad \square$$

Two consequences. First, an entry reads $+1$ only when $p_i = p_j$, so a rare class and a common class have a lower ceiling than two common classes, and entries for pairs with different inclusion rates are not on a common scale. Second, as α grows the sets shrink, every p_i falls, and every floor entry (3) moves toward zero whether or not the classifier changed. The paper sweeps α over four values and presents the dependence as a feature.

Lemma 2 (Identification of the inclusion rates). *Suppose every prediction set is a singleton and at least three classes have inclusion rate strictly between zero and one. Then for any three distinct such classes i, j, k ,*

$$o_i = -\frac{\rho_{ij} \rho_{ik}}{\rho_{jk}}, \quad p_i = \frac{o_i}{1 + o_i}. \quad (4)$$

The off-diagonal entries of the matrix therefore determine every inclusion rate exactly.

Proof. With singleton sets no two classes ever share a set, so (3) holds for every pair, and $\rho_{ij}\rho_{ik} = \sqrt{o_i o_j} \sqrt{o_i o_k} = o_i \sqrt{o_j o_k} = -o_i \rho_{jk}$. Divide by $-\rho_{jk}$, which is nonzero since $o_j, o_k > 0$. The map $o \mapsto o/(1 + o)$ inverts $p \mapsto p/(1 - p)$. $\square$

The paper argues that the matrix is privacy preserving because it “only contains probabilities” and so “does not leak information on the size and composition” of a client’s data. Correlations can be negative, so the matrix does not contain probabilities. A row normalized confusion matrix does, and reveals no counts either. And in the singleton regime, which is the regime of an accurate classifier at moderate α , Lemma 2 recovers the inclusion rates, which for an accurate classifier are the predicted class frequencies, which approximate the label histogram. The matrix is offered as useful precisely because it exposes how a client’s data differs from the others. No threat model, mechanism or leakage bound is given.

5 What the coverage guarantee constrains

The paper states that “for a given input sample” the set “is guaranteed to contain the correct class” with probability close to $1 - \alpha$. The guarantee (2) is marginal, not conditional on the input, and distribution free conditional coverage is impossible in finite samples (Barber et al., 2021). The distinction matters here because the authors go on to interpret individual class pairs and individual federated clients as if validity transferred to them. It does not transfer to a client or a class either. Take ten equally weighted clients, nine with coverage one and one with coverage zero. Pooled coverage is exactly 0.9, which meets a nominal $1 - \alpha = 0.9$, and the tenth client has no protection at all. The same arithmetic applies to a class with a tenth of the data that is missed every time.

The paper also writes the upper coverage term as $1/(n - 1)$, which is looser than the standard $1/(n + 1)$ (Lei et al., 2018; Stutz et al., 2022) and probably a typographical error. And the federated experiments compare local matrices computed against a global calibration set across clients that are non identically distributed by construction. Under that mismatch even the marginal guarantee need not hold, and no coverage or set sizes are reported for any client.

6 What to report instead

If the question is which classes the classifier confuses, the probabilities answer it directly. The soft confusion matrix

$$C_{ab} = \mathbb{E}[p_b(X) \mid Y = a]$$

is the average probability assigned to class b on inputs of true class a . It is directional, it sees near misses at every resolution, and it needs no calibration set and no threshold. It needs labeled evaluation data, as any confusion matrix does. If labels are unavailable, the mean pairwise product $\mathbb{E}[p_i(X)p_j(X)]$ is symmetric and label free. If sets are genuinely the object, because a downstream consumer receives sets and nothing else, then the co-occurrence rate q_{ij} is the quantity the abstract describes, and the coverage confusion matrix of Stutz et al. (2022), which conditions on the true

label, is directional. In every case the inclusion rates and the set sizes should be reported alongside, since by Lemma 1 an entry of (1) cannot be read without them. The paper compares against none of these. Its only comparison is with the hard confusion matrix, which discards the probabilities entirely, so the contrast it draws is soft versus hard rather than conformal versus not.

The paper’s client selection experiment does not bear on any of this, but we note that it selects the clients with the highest entries for two critical classes while the centralized section reads a falling entry as the classifier learning to separate them, that its Figures 6 and 7 are captioned as matrix entries while their axes are class accuracy, and that its Table 1 reports overall CIFAR-100 accuracy falling from 0.511 to 0.493 under a heading of no effect, without repetitions, seeds or error bars.

Simulations that exhibit Examples 4, 6 and the direction blindness of Example 1 on a synthetic ten class classifier with adaptive prediction sets, and that compare the matrix with C_{ab} and $\mathbb{E}[p_i p_j]$ as α varies, are available at <https://conformalprediction.net/demos/32-conformal-correlation.html>.

7 Conclusion

The conformal correlation matrix is the phi coefficient of a thresholded softmax. By Proposition 1 it sees only the law of the sets, and by Proposition 2 the coverage guarantee leaves every entry free. The seven examples show that each reading the authors attach to an entry, positive for similar, negative for dissimilar, zero for unrelated, fails on a three class problem with perfect coverage. Lemma 1 shows that the entries the authors read as dissimilarity are functions of the inclusion rates alone, and Lemma 2 shows that in the regime of an accurate classifier those rates can be read off the matrix exactly. None of this is repaired by the certificate, which concerns the true label and says nothing about which other labels sit beside it.

References

- R. F. Barber, E. J. Candès, A. Ramdas, and R. J. Tibshirani. The limits of distribution-free conditional predictive inference. *Information and Inference*, 10(2):455–482, 2021.
- J. Lei, M. G’Sell, A. Rinaldo, R. J. Tibshirani, and L. Wasserman. Distribution-free predictive inference for regression. *Journal of the American Statistical Association*, 113(523):1094–1111, 2018.
- A. Perlo, C. F. Chiasserini, G. De Veciana, and F. Malandrino. Characterizing the performance of classification models through conformal correlation matrices. *Computer Communications*, 247:108398, 2026.
- Y. Romano, M. Sesia, and E. J. Candès. Classification with valid and adaptive coverage. In *Advances in Neural Information Processing Systems*, volume 33, pages 3581–3591, 2020.
- D. Stutz, K. Dvijotham, A. T. Cemgil, and A. Doucet. Learning optimal conformal classifiers. In *International Conference on Learning Representations*, 2022.
- V. Vovk, A. Gammerman, and G. Shafer. *Algorithmic Learning in a Random World*. Springer, 2005.
- G. U. Yule. On the methods of measuring association between two attributes. *Journal of the Royal Statistical Society*, 75(6):579–652, 1912.