

# Conformal Prediction as a Transform within a Grammar for Probabilistic Forecasting

Peter Cotton  
peter.cotton@microprediction.com

August 2026

## Abstract

Conformal procedures factor into a model, a scalar residual, and a rank-calibration operator. We study the fence-post rank map as a stateful operator in a grammar of transforms. Interpolated, it is invertible and, on distinct exchangeable input with its slope guard inactive, exact at retained levels, within one knot cell everywhere. The certificate attaches to the operator’s immediate output and is not inherited automatically. On economic series a serial model after the operator outscores the literal stop on average, while against the same model without it, CRPS improves and log score worsens. A Bayes pool admits it at bounded cost.

## 1 Introduction

In the canonical residual-pooling pattern, a forecaster fits a conditional model and then quotes uncertainty by ranking its past errors with every state pooled. The ranking is honest on average and silent about tonight. That tension, conditional effort upstream and unconditional treatment at the end, is the subject of this paper.

Split conformal prediction can be written as a chain. Fit a model, take residuals, apply the fence-post empirical map

$$E_n(r) = \frac{1 + \#\{i \leq n : R_i \leq r\}}{n + 1}, \quad (1)$$

and map the resulting rank back through the model to a prediction set (Vovk et al., 2005; Lei et al., 2018). The value of (1) at a new point is a conformal p-value, and the map itself is the unrandomized conformal predictive distribution of Vovk et al. (2019).

We call it the fence-post transform, after the two +1 corrections:  $n$  points cut the line into  $n + 1$  cells, and the guarantee rides on the count of cells, not of points. The knot probabilities  $i/(n + 1)$  are Weibull’s plotting positions, the means of uniform order statistics, so the object is classical.

The literature mostly treats this chain as a construction to be studied whole, and its many variants, normalized, Mondrian, quantile, distributional, and temporal (Boström et al., 2021; Romano et al., 2019; Chernozhukov et al., 2021; Gibbs and Candès, 2021; Auer et al., 2023), as separate constructions. The nearest exceptions stop one step short. Conformal predictive systems made the map emit a distribution function, and conformal calibrators (Vovk et al., 2020) apply that map to another system’s output, one level of composition with the map still the last word. We propose a different organization, in which the map’s position is a variable. The fence-post transform is one operator. The chain is one production rule in a grammar of

composable transforms. The interesting question is then not which conformal method to choose but which sequence of transforms minimizes the forecaster’s scoring rule.

Stated plainly, conformal prediction is not, in general, an alternative to statistical modeling. It is a particular rank-based transform of a model’s output. Its finite-sample property distinguishes it from fitted recalibrators at the point where it is applied. It does not privilege stopping there, it does not remove the modeling that follows, and it does not determine predictive quality under a proper scoring rule.

The standard conformal pattern also carries a tension the present work avoids. A skeptic of the standard conformal prediction patterns might reasonably ask why an approach that essentially adopts the empirical distribution, in a sense advocating not modeling at all, is appropriate in the context of model residuals but not in general (for instance, as a replacement for the base model).

The stock answer is about targets. For marginal coverage, pooling is justified by exchangeability of the residuals alone, a weaker and testable hypothesis, and distribution-free conditional coverage is provably out of reach (Vovk, 2012; Lei and Wasserman, 2014; Foygel Barber et al., 2021), so the pooled law is what one settles for. That answer does not carry over to a forecaster scored by a proper scoring rule. There the pooled residual law is quoted as a predictive density, its adequacy is exactly conditional invariance, and the certificate does not establish that. Certifying ranks under exchangeability makes stopping at the pooled law a modeling decision, not a consequence of the guarantee, and pooling the observable itself, the limiting case of an inadequate model, would frequently be ludicrous.

In univariate probabilistic forecasting, Platt scaling, isotonic and quantile recalibration, temperature scaling, conformal predictive systems, and distributional conformal prediction can all be represented as monotone recalibration operators on a scalar predictive, score, or probability coordinate, differing in what they fit and what they guarantee (Section 5). The fence-post transform is the member with a certificate, and the certificate is a positional type annotation on one operator inside a composable predictive program: a reason to admit the operator to the grammar, and nothing more.

This reorganization has content because the fence-post operator composes. Once interpolated,  $E_n$  is an invertible change of variables with a density, and its probit composition  $z = \Phi^{-1}(E_n(r))$  maps exchangeable residuals to approximately standard Gaussian coordinates. Nothing obliges a forecaster to stop there. One can fit further structure on the Gaussianized stream, transform again, and continue, exactly as one does with any other member of the transform library. Conformal calibration stops at the transform. The grammar keeps going.

Two testable claims follow. The first is that the operator is not a natural stopping point: a serial model fitted on the Gaussianized stream can beat quoting the transformed empirical law where it stands, when exploitable residual structure remains. The second is that nesting is cheap: a pool over chains that includes the conformal pattern should not pay materially for admitting it, because the pool reallocates weight away from the pattern on series where it is a bad choice. We confirm both on a corpus of economic series and quantify a third finding we did not anticipate: the operator’s incremental value inside a fixed chain depends on the scoring rule, and the dependence is not explained by the smoothing we tested.

A companion paper prices what a fixed representation choice costs (Cotton, 2026b). Under logarithmic loss the irreducible cost of pooling residuals in given coordinates is a mutual information, the conformal information gap, and the companion measures it on the same corpus. The present paper is the constructive complement. If a fixed conformal pattern can be a bad choice

with a quantifiable price, the remedy is not a better fixed choice. It is a grammar in which the choice is made by the scoring rule, online, per series.

## 2 The fence-post transform as an operator

We work with online univariate transforms in the sense of the `skaters` library, described at length in a methods paper (Cotton, 2026d) and briefly in a software paper (Cotton, 2026e). A transform is a pair of functions. The forward function consumes one observation and its own past state and emits a transformed observation. The inverse function maps predictive distributions in the transformed space back to the original space. Conditional on its own past state, the forward map is an invertible change of variables in the current observation, so predictive densities push back exactly, Jacobians included.

The raw fence-post map (1) is a step function. It has no density and no inverse, so it is not a transform in this sense. The minimal repair is interpolation. Let  $y_{(1)} \leq \dots \leq y_{(n)}$  be the past observations, place knots at the order statistics with fence-post probabilities  $p_{(i)} = i/(n+1)$  for  $i = 1, \dots, n$ , set  $z_{(i)} = \Phi^{-1}(p_{(i)})$ , and interpolate the pairs  $(y_{(i)}, z_{(i)})$  piecewise linearly. We call the resulting operator *gaussianize*. The knot values are the left limits of the right-continuous map (1) at the calibration points, and this convention is what makes the certificate exact at the knots (Proposition 3). That is the full map. The implementation retains at most 39 of the knots, chosen at equal mass from the full sorted history so that each retained knot is an exact order statistic until the slope guard below binds, and drops any knot whose value ties its neighbor, so ties widen the retained cells. We write *gaussianize-full* for the map through every knot and *gaussianize* for the implemented retained-knot member. Proposition 5 prices the difference: the certificate is a function of the retained resolution.

The interpolation changes the object, and Section 3 states exactly what survives the change: the discrete map’s output is a function of ranks alone, while the interpolated map’s output depends on spacings as well, so the two are different operators with different certificates.

Beyond the edge knots the map extends linearly in  $(y, z)$  coordinates, and linearity in the probit coordinate is a Gaussian tail in the observable, with scale set by the spacing of the extreme quantiles. Tail mass is unobservable, so any usable version of the operator carries a tail assumption, and this is the mildest one consistent with the coordinates. A kernel-smoothed variant replaces the empirical distribution function with a Gaussian-kernel estimate at Silverman bandwidth, represented by a monotone  $C^1$  cubic through the same knots, and we use it in Section 4.4 to test whether smoothness matters.

Two guards complete the definition, and both were forced by data rather than taste. First, every knot segment’s inverse slope  $dy/dz$  is capped at twenty times the interquartile inverse slope, so the forward slope is bounded below. An upstream variance collapse can plant a many-sigma outlier as an edge knot, and an uncapped edge segment then hands the inverse map an explosive extrapolation slope. The cap is an explicit assumption: tail scale at most a fixed multiple of the central scale.

Second, the value handed downstream is clamped to  $|z| \leq 8$ . The clamp applies only to the coordinate supplied to downstream state updates and is not part of the invertible map used to construct or pull back predictive densities, so the change-of-variables identity is intact. When it binds, the downstream learner is updated with a bounded surrogate rather than the literal transformed observation, and that alteration can persist through the learner’s state. In the type system the two objects are distinct: *gaussianize* is invertible and density-carrying, and the

bounded update channel is a winsorized learner input, not a transform.

Within the knot range the clamp cannot bind, since no fence-post rank exceeds probability  $n/(n+1)$  and  $\Phi^{-1}(n/(n+1)) < 8$  for any feasible  $n$ . It can bind on the extended tail and on degenerate warmup emissions, and the latter is why it exists, because downstream learners have expanding memory. In one series in our corpus a degenerate warmup emitted a single transformed value near 8000, and the downstream Welford variance and recursive least-squares state never forgot it, inflating that series' loss by four orders of magnitude. Composability is not free. An operator that is harmless as a terminal step can poison downstream learners through one bad emission, and the operator's definition has to prevent this, not the caller.

The guards are better read as declared budgets than as patches. Invertibility alone is too weak a property for a chain. The pullback quantile  $y_\alpha = m^{-1}(q(\alpha))$  inherits the steepness of the inverse map  $m^{-1}$ , and beyond the edge knots the inverse is an assumption rather than an estimate, because tail mass is unobservable at any  $n$ . The slope cap makes the assumption finite. It declares that  $m^{-1}$  is Lipschitz, with constant  $L$  equal to twenty times the central inverse slope, and the declaration buys a guarantee.

The declaration also has a price, and three objects need separating. The exact-knot interpolant is what the propositions of Section 3 certify. The slope-guarded map is what the experiments run: the cap is enforced by lowering any knot whose segment is too steep, so when it binds, the map no longer passes through the displaced order statistics, and retained-level exactness is suspended there. The clamped emission is a third object, a winsorized channel to downstream state and not part of the invertible map. Remark 7 prices the middle one.

**Lemma 1** (Pullback quantiles are budgeted). *Let  $m$  be a strictly increasing forward map whose inverse is  $L$ -Lipschitz on all of  $\mathbb{R}$ , extrapolation included, and let  $q$  be the quantile function of the predictive in transformed coordinates. Then the pulled-back predictive has quantiles  $y_\alpha = m^{-1}(q(\alpha))$  and*

$$|y_\alpha - y_{1/2}| \leq L |q(\alpha) - q(\tfrac{1}{2})| \quad \text{for every } \alpha \in (0, 1).$$

*For a chain of such maps the same bound holds with  $L$  the product of the rungs' constants.*

*Proof.* Monotone maps transport quantiles, so  $y_\alpha = m^{-1}(q(\alpha))$ . The Lipschitz bound applied to the pair  $q(\alpha), q(\frac{1}{2})$  gives the display. A composition of Lipschitz maps is Lipschitz with the product constant.  $\square$

With a Gaussian leaf of scale  $\sigma$  the right side is  $L\sigma |\Phi^{-1}(\alpha)|$ . The budget is conditional on the operator's state:  $L$  is twenty times the current central inverse slope and  $\sigma$  is fitted, so the lemma bounds quoted quantiles relative to the current central scale, not by a universal constant. Within that reading, no realization amplifies a quoted quantile past its budget. The global  $L$  is also purchased with the knot displacement of Remark 7: the Lipschitz budget and retained-level exactness are competing constraints, and no operator holds both exactly on arbitrary data. The naive empirical map declares  $L = \infty$  and buys nothing. The clamp is the same discipline in the forward direction: the codomain bound  $|z| \leq 8$  is intrinsic to the fence-post ranks, and a downstream learner with expanding memory is safe against exactly the emissions the bound excludes. Any finite multiplier yields the lemma, so the value twenty is the one remaining choice of taste, a tail-conservatism dial. Fitting the tail where the sample reaches, generalized-Pareto by censored maximum likelihood as the leaf already does, with the budget kept as a prior bound, is the refinement we leave open.

Before enough observations accumulate, the operator standardizes by expanding moments, which is the  $n \rightarrow 0$  limit of the empirical map. After the warmup the knots are rebuilt each

step, at most thirty-nine of them at equal mass from the full sorted history, so the retained knots sit at exact order statistics, and carry no rank-approximation error while the slope guard is inactive. The sorted buffer itself grows with the series. A bounded compacted sketch is the obvious refinement, and its price is the rank-error term of Proposition 5.

### 3 What composition preserves

The grammar can be said minimally and formally. A *chain* is a finite word  $C = (T_1, \dots, T_m, L)$ , stateful transforms followed by a terminal leaf, and a transform is a tuple: a forward emission, an inverse pullback, a state update, a clock, input and output types, and declared properties. The forward map consumes an observation and the operator's state and emits a coordinate. The pullback carries predictive laws backward. The clock says when the state may absorb data. The properties are judgments, monotone, invertible, density-carrying, bi-Lipschitz with a declared budget, exchangeability-preserving, previsible, or  $\delta$ -calibrated at a position.

Composition obeys checkable rules. Invertible composed with invertible is invertible, and Lipschitz budgets multiply (Lemma 1). A fixed strictly increasing link transports the grid certificate, with its levels transported along the link, because ranks are invariant under it. Previsible composed with previsible is previsible, since the composite map at time  $t$  is built from maps measurable before  $t$ . A calibration certificate is not, in general, inherited through a fitted state-dependent operator, which is the second fact below. A pool over chains is itself an operator on chains, and Proposition 9 is its admission rule.

Composition preserves some properties and destroys others, and the pattern is positional. We record the four facts that organize our experiments.

**Proposition 2** (Fence-post uniformity). *Let  $R_1, \dots, R_{n+1}$  be exchangeable and almost surely distinct. Then  $E_n(R_{n+1})$  is uniform on the grid  $\{k/(n+1) : k = 1, \dots, n+1\}$ . In particular,  $\Pr\{E_n(R_{n+1}) \leq k/(n+1)\} = k/(n+1)$  for  $k = 1, \dots, n+1$ .*

*Proof.* By exchangeability every ordering of  $R_1, \dots, R_{n+1}$  is equally likely, and by almost-sure distinctness the rank of  $R_{n+1}$  among the  $n+1$  values is uniform on  $\{1, \dots, n+1\}$ . The map (1) evaluated at  $R_{n+1}$  equals the rank divided by  $n+1$ .  $\square$

This is the standard conformal fact (Vovk et al., 2005) stated as a property of the operator's output. It certifies the discrete map (1), and the discrete map is what the experiments do not run. The interpolated operator of Section 2 uses spacings as well as ranks, and its output between knots is not distribution-free. What survives interpolation is the grid.

**Proposition 3** (Grid calibration survives interpolation). *Let  $Y_1, \dots, Y_{n+1}$  be exchangeable and almost surely distinct. Let  $m$  be any strictly increasing map, interpolation and tails arbitrary, that passes through the knots:  $m(y_{(i)}) = z_{(i)}$  at the order statistics  $y_{(1)} < \dots < y_{(n)}$  of  $Y_1, \dots, Y_n$ . Then*

$$\Pr\{m(Y_{n+1}) \leq z_{(i)}\} = \frac{i}{n+1}, \quad i = 1, \dots, n.$$

*Proof.* Monotonicity gives  $\{m(Y_{n+1}) \leq z_{(i)}\} = \{Y_{n+1} \leq y_{(i)}\}$ . By exchangeability and distinctness the rank of  $Y_{n+1}$  among all  $n+1$  values is uniform on  $\{1, \dots, n+1\}$ , and  $Y_{n+1} \leq y_{(i)}$  holds for exactly the  $i$  lowest ranks.  $\square$

There is no need to insist on exactness off the grid, because the relaxation is uniformly small and assumption-free. Call an operator  $\delta$ -conformal if, on exchangeable almost surely distinct input, its output  $U$  satisfies  $\sup_u |\Pr\{U \leq u\} - u| \leq \delta$ .

**Proposition 4** (Interpolation is  $1/(n+1)$ -conformal). *In the setting of Proposition 3, write  $U = \Phi(m(Y_{n+1}))$  for the probability-scale output of any interpolated map through the knots. Then*

$$\sup_{u \in (0,1)} |\Pr\{U \leq u\} - u| \leq \frac{1}{n+1},$$

*whatever the interpolation, the tails, and the continuous input law. Equivalently,  $m(Y_{n+1})$  is within Kolmogorov distance  $1/(n+1)$  of standard Gaussian.*

*Proof.* For  $u$  between consecutive grid levels  $i/(n+1)$  and  $(i+1)/(n+1)$ , monotonicity of the distribution function and Proposition 3 trap  $\Pr\{U \leq u\}$  in  $[i/(n+1), (i+1)/(n+1)]$ , the same cell that contains  $u$ . Two points of one cell differ by at most  $1/(n+1)$ . The edge cells are the same argument with endpoints 0 and 1.  $\square$

The bound generalizes to the implemented operator, which does not keep every knot.

**Proposition 5** (The certificate scales with the retained resolution). *Let  $0 = r_0 < r_1 < \dots < r_K < r_{K+1} = n+1$ , and let  $m$  be any strictly increasing map with  $m(y_{(r_j)}) = \Phi^{-1}(r_j/(n+1))$  at the retained order statistics, interpolation and tails arbitrary. Then  $U = \Phi(m(Y_{n+1}))$  satisfies*

$$\sup_{u \in (0,1)} |\Pr\{U \leq u\} - u| \leq \Delta_n := \max_{0 \leq j \leq K} \frac{r_{j+1} - r_j}{n+1}.$$

*If each retained knot's value sits within  $e$  stream ranks of its nominal rank, the bound degrades to  $\Delta_n + e/(n+1)$ . The hypothesis of Proposition 3, exchangeable and almost surely distinct input, is in force.*

*Proof.* The argument of Proposition 3 gives exactness at each retained level  $r_j/(n+1)$ . Between consecutive retained levels, monotonicity traps  $\Pr\{U \leq u\}$  and  $u$  in one retained cell, of width at most  $\Delta_n$ . For the rank-error term, the hypothesis is a bracketing contract: each interpolation knot  $\bar{y}_j$  satisfies  $y_{(r_j-e)} \leq \bar{y}_j \leq y_{(r_j+e)}$  almost surely. Event inclusion then traps  $\Pr\{Y_{n+1} \leq \bar{y}_j\}$  between  $(r_j - e)/(n+1)$  and  $(r_j + e)/(n+1)$ , so the level error at a knot is at most  $e/(n+1)$ , which adds to the cell width.  $\square$

The full map is the case  $\Delta_n = 1/(n+1)$ . For the implemented operator the accounting has two regimes. While the slope guard is inactive the retained knots are exact order statistics,  $e = 0$ , and after warmup the certificate is roughly one part in forty everywhere, exact at the retained levels. When the guard binds it lowers the offending knots, the map no longer passes through the displaced order statistics, and the displacement enters  $e$  through the bracketing contract. The displacement need not be small, as the next remark shows, so the guarded operator is not  $\Delta_n$ -conformal unconditionally. The  $e$  term is also the price a future bounded-memory compaction would pay. Ties need their own statement.

**Corollary 6** (Tie handling). *Drop the distinctness assumption. For the discrete rank, randomization restores exactness: ordering tied observations uniformly at random makes the rank of  $Y_{n+1}$  uniform on  $\{1, \dots, n+1\}$  whatever the tie structure, so Proposition 2 survives ties unchanged. Interpolation admits no such fix, because tied values are numerically equal and reordering them*

does not change the map. For the interpolated operator with one knot per distinct retained value, let  $\beta_n$  be the largest tie-block mass and let  $\Delta_n$  be the largest retained-cell width, maximized over which slot of the multiset the new observation occupies. Then  $\sup_u |\Pr\{U \leq u\} - u| \leq \Delta_n + \beta_n$ , conditionally on the multiset of the  $n + 1$  values and hence unconditionally, and exactness at retained levels is lost.

*Proof.* Given the multiset, the slot of  $Y_{n+1}$  among the  $n + 1$  values is uniform, by exchangeability. The knots are built from the remaining  $n$  values, so they vary with the slot, which is why  $\Delta_n$  is maximized over slots: so defined it is a function of the multiset alone. For a retained knot at the value  $y$  with level  $\ell$ ,  $\Pr\{Y_{n+1} \leq y \mid \text{multiset}\}$  is the mass of tie blocks starting at or below  $y$ , and it differs from  $\ell$  by at most the mass of the block at  $y$  together with the one-slot shift of the level itself, in total at most  $\beta_n$ . The sandwich argument of Proposition 5 then runs conditionally with every cell widened by  $\beta_n$ , and averaging over the multiset gives the unconditional bound.  $\square$

*Remark 7* (The guard is a third price). Take  $n = 78$  calibration values, seventy-one at  $0, 1, \dots, 70$  and eight near  $10^9$ , all distinct, in exchangeable order. With 39 retained knots the top knot is the seventy-eighth order statistic and  $\Delta_n = 2/79$ . The slope cap, at twenty times the central inverse slope, drags that knot from the high cluster down to roughly 671, between the clusters. At the retained level 78/79 the probability that a new exchangeable draw falls below the displaced knot is 71/79, an error of 7/79, exceeding  $\Delta_n$ . Retained resolution and tie mass are therefore not the whole price of the certificate. Guard displacement is a third term, zero exactly while the guard sleeps. Deleting offending knots instead of moving them restores exact order statistics at an enlarged  $\Delta_n$ , and the operator implements it as an option. On the campaign corpus the deletion guard is indistinguishable from the moving cap, identical on all but three percent of series with a mean difference under 0.0003 nats, so the unconditional certificate costs nothing measurable.

The corpus of Section 4 filters to repeat fraction below five percent rather than zero, so  $\beta_n$  is small and not vanishing there. The statement worth keeping is that the certificate is a function of knot resolution, sketch error, and tie mass: the guarantee is priced in the operator's computational budget.

The discrete fence-post map and gaussianize-full both have global error at most  $1/(n + 1)$ : the discrete map's output is supported on the grid, so as a distribution on  $[0, 1]$  it too sits within Kolmogorov distance  $1/(n + 1)$  of uniform, and only the full map is invertible. The implemented retained-knot map has global error at most  $\Delta_n$ , exact at retained levels under distinctness.

What separates them is the exact statement. The discrete map is deterministic, exactly calibrated at every grid level, noninvertible, and  $1/(n + 1)$ -calibrated between grid levels. Randomized smoothing, which spreads each output uniformly over its rank cell, is exactly calibrated at every level and is a stochastic kernel rather than a deterministic transform (Vovk et al., 2005). The interpolated map is deterministic and invertible, stays exact on the grid, and incurs at most one full-grid cell of calibration error elsewhere. The grammar uses the deterministic, invertible version and quantifies its one-cell calibration relaxation.

The exact conformal object is a rank certificate, and a rank certificate is not a predictive density. Turning it into one requires interpolation, smoothing, tail assumptions, and model selection, the same choices met everywhere else in probabilistic forecasting, and exchangeability supplies none of them.

The log score will still notice: it integrates the predictive density between grid points and in the tails, exactly where the certificate is  $\delta$ -relaxed and assumption-laden, and Section 4.4 measures what that costs.

The second fact is that this guarantee is positional. Proposition 2 concerns the operator’s immediate output. If a downstream model is fitted on  $z_1, \dots, z_t$  and its parameters enter later forecasts, the composed system’s forecast at time  $t + 1$  depends on the ordering of the past stream, and the exchangeability that delivered uniformity no longer applies to the composed output. The calibration certificate attaches to a position in the chain, not to the chain. This is not a defect. It is the familiar price of fitting, and the reason the companion paper’s ledger separates representation, model-form, and estimation cost (Cotton, 2026b). But it means the grammar has a typing discipline: an operator’s guarantee names the position at which it holds, and claims about the composed system need their own arguments.

The third fact concerns schedules, and it explains what the conformal literature’s data splitting is for and how an online chain does without it. Call a chain *previsible* if the map applied to each observation is measurable with respect to information strictly prior to that observation. Dual use, the defect the splitting exists to avoid, is exactly a previsibility violation: an observation flowing through a map whose state already contains it.

Split conformal enforces previsibility with a coarse clock. Interior operators learn on a fitting block and freeze, the terminal fence-post learns on a disjoint calibration block and freezes, and test points meet only frozen maps.

The chains of this paper enforce the same rule with the finest clock. Every operator emits before it absorbs, so the map applied at time  $t$  is built from  $y_1, \dots, y_{t-1}$  alone, and the past is a calibration set that never contains the present. Batching is one scheduler for one rule, and updating after emission is another.

On the exchangeable hypothesis the fine clock matches the batch guarantee and adds independence across time.

**Proposition 8** (Sequential fence-post outputs are independent). *Let  $Y_1, \dots, Y_T$  be exchangeable and almost surely distinct, and let  $U_t = E_{t-1}(Y_t)$  be the fence-post transform of  $Y_t$  through the map built from  $Y_1, \dots, Y_{t-1}$ , as in (1) with  $n = t - 1$ . Then  $U_1, \dots, U_T$  are independent, and  $U_t$  is uniform on  $\{1/t, \dots, t/t\}$ .*

*Proof.* Let  $K_t \in \{1, \dots, t\}$  be the rank of  $Y_t$  among  $Y_1, \dots, Y_t$ . From (1) with  $n = t - 1$ ,  $U_t = (1 + \#\{i < t : Y_i \leq Y_t\})/t = K_t/t$ . The full ranking of  $Y_1, \dots, Y_T$  is a permutation  $\pi$  of  $\{1, \dots, T\}$ , and the map  $\pi \mapsto (K_1, \dots, K_T)$  is a bijection onto  $\{1\} \times \{1, 2\} \times \dots \times \{1, \dots, T\}$ : given the sequential ranks, the permutation is reconstructed by inserting each element at its stated position, and both sets have cardinality  $T!$ . Exchangeability with almost-sure distinctness makes  $\pi$  uniform over the  $T!$  orderings, so  $(K_1, \dots, K_T)$  is uniform over the product set. Uniformity over a product set is independence of the coordinates, each uniform on its factor.  $\square$

The independence of sequential ranks is classical (Rényi, 1962) and powers online exchangeability testing (Vovk et al., 2003). The proposition concerns the discrete sequential ranks alone. The interpolated coordinates carry spacing information, and no independence is claimed for them.

The narrowing matters, because previsibility is not a substitute for splitting. Previsibility prevents leakage, the map at time  $t$  not containing  $y_t$ , and that is all it prevents. It does not make the transformed stream exchangeable: previsible strictly increasing maps that differ across time destroy exchangeability trivially, as when every second observation is shifted by a constant. Splitting does more. It freezes one map for calibration and test, so their scores stay exchangeable whenever the data are. The accurate statement is that emit-before-absorb guarantees prequential out-of-sample evaluation, and if the stream arriving at the fence-post

rung is itself exchangeable, its sequential ranks are independent. The chains of Section 4 do not satisfy that hypothesis, and their certificate reads through the idealization already declared. Here the proposition separates the schedules. Split conformal’s p-values share one calibration set and are dependent. The prequential outputs are independent, and the price is early resolution: the map at time  $t$  has granularity  $1/t$ , hence the warmup of Section 2.

The certificate survives any further fixed monotone map, the probit included, because monotone maps preserve ranks. It is not automatically inherited through a downstream fit, by the second fact, and a chain that wants a terminal certificate ends with one more fence-post step on the final residuals. Full conformal is a third schedule, symmetrizing the whole chain over the data at the cost of refitting per candidate value (Vovk et al., 2005). We do not use it. An operator’s annotation is therefore a position and a clock.

For the economic series of Section 4 the exchangeable hypothesis of Proposition 8 is false, as it is throughout time-series conformal prediction, so the uniformity claims are conditional on model adequacy. Previsibility is not. It holds by construction, and every score in this paper is out-of-sample even where the guarantee hypothesis fails. The contracts so far assemble into the beginning of the type table the conclusion asks for:

| operator                  | previsible        | exch.-preserving  | invertible  | certificate                                                                                   |
|---------------------------|-------------------|-------------------|-------------|-----------------------------------------------------------------------------------------------|
| fence-post, discrete      | yes               | indep. seq. ranks | no          | exact at grid levels                                                                          |
| gaussianize, full grid    | yes               | not claimed       | yes         | exact on all $n$ rank levels, $1/(n+1)$ globally                                              |
| gaussianize, implemented  | yes               | not claimed       | yes         | exact on retained levels (distinct input, guard inactive), $\Delta_n$ plus guard displacement |
| fixed monotone link       | yes               | yes               | yes         | transparent                                                                                   |
| EMA location, EWMA scale  | yes               | no                | given state | none                                                                                          |
| frozen split-fitted model | yes, coarse clock | cal/test pairs    | varies      | none                                                                                          |
| refit on the stream       | yes               | no                | given state | none, blocks pass-through                                                                     |

The middle column is the working hypothesis of every chain in Section 4: location and scale do not preserve exchangeability, so the fence-post rung downstream of them operates under an idealization, which the contracts state rather than hide.

The fourth fact prices the choice of coordinates. The companion paper’s transform-pooling theorem states that pooling residuals  $R$  after an invertible state-dependent change of coordinates  $\phi_X$  has irreducible log-score gap  $I(\phi_X(R); X)$ , the mutual information between the transformed residual and the state  $X$  on which its true conditional law depends: the Jacobians of the transform cancel, and only the coordinates’ ability to remove conditional heterogeneity matters (Cotton, 2026b). The gaussianize operator is a data-dependent monotone map, measurable with respect to its own past state, and piecewise  $C^1$ , so conditional on the state the theorem’s hypotheses hold. The grammar view adds only this: since operators compose, the coordinates in which one pools are themselves a product of the chain, and the gap is a property of the whole chain up to the pooling point.

The operator’s ancestry is long. The branches are not one idea. They transform different objects for different reasons, and what they share is the trunk: push an observation through a distribution function and work on the flat scale.

The empirical distribution function is Galton’s ogive (Galton, 1875). Plotting positions place ordered observations at probabilities of the form  $i/(n+1)$  (Hazen, 1914), and the probability

integral transformation received its name and first systematic treatment from Pearson (1938). Normal scores, the replacement of ranks by standard-normal quantiles, are the gaussianize map’s direct ancestor (van der Waerden, 1952). That estimation spoils the transform’s exactness has been known since David and Johnson (1948), and the positional fact above is its descendant.

Rosenblatt’s chaining through conditional distributions whitens any process (Rosenblatt, 1952), and normalizing flows are its learned, offline descendants (Tabak and Vanden-Eijnden, 2010; Papamakarios et al., 2021). Parametric monotone Gaussianizers run from Box and Cox (1964) to Yeo and Johnson (2000), quantile mapping in the geosciences is the same move under another name (Panofsky and Brier, 1968), transformed residuals as model criticism begin with Cox and Snell (1968), the PIT as a forecast-evaluation instrument with Diebold et al. (1998), and the prequential stance that all of this paper’s scoring obeys is Dawid (1984).<sup>1</sup>

The transform also ran as a market before it ran here as an operator. The microprediction platform’s z-streams solicited competitive real-time predictions of PIT-transformed residual streams (Cotton, 2022), chaining solicitations so that one crowd predicted the residuals of another’s forecasts. That is interior composition as a mechanism, and the algebra of such mechanisms is developed in Cotton (2026a) and Cotton (2026c). The grammar of this paper is the single-forecaster form of the same design.

Against this ancestry the fence-post operator’s distinction is narrow and real: at its position, on exchangeable input, its ranks carry a finite-sample certificate.

## 4 Measurements

### 4.1 Design

The corpus, scoring conventions, and evaluation protocol follow the companion paper (Cotton, 2026b). We use 701 FRED change series with at least 600 observations and repeat fraction below five percent. The series are latest-vintage changes, so the claims concern sequential modeling of economic-shaped sequences, not operational real-time forecasting, where revisions and publication lags would enter. Every forecaster runs prequentially: at each time it issues a one-step predictive density from data strictly before that time, then sees the observation and updates. Scoring starts after a burn-in of 300 observations. The predictive in the original coordinates is the finite Gaussian mixture produced by the nine-particle inverse. That implemented mixture, blended with weight 0.01 toward an expanding-window Gaussian reference, is scored exactly, with no clipping or flooring in the reported nats. CRPS is computed in closed form from the Gaussian-mixture predictive. Per-series CRPS is normalized by the location baseline defined next, so scales are comparable across series.

The bare leaf, six single chains, and three pools are compared. All chains end in a Gaussian leaf with online variance, and all transforms update online.

---

<sup>1</sup>A fuller annotated timeline is maintained at <https://skaters.microprediction.org/heritage.html>.




Figure 1: The five comparisons behind the three findings, under both scoring rules, with each duel’s corpus in its label. Bars to the right favor the first-named config. Log score improvement is in nats per observation. CRPS improvement is the mean per-series ratio, expressed as the percent cut. Win rates are per-series.

| config | chain                                                                          |
|--------|--------------------------------------------------------------------------------|
| A      | location (EMA) → leaf                                                          |
| B      | location → scale (EWMA standardize) → leaf                                     |
| C      | location → gaussianize → leaf (rank-transform chain)                           |
| D      | location → scale → gaussianize → leaf (normalized rank-transform chain)        |
| E      | location → scale → gaussianize → AR(1) → leaf (refit through)                  |
| F      | location → scale → AR(1) → leaf (same refit, no gaussianize)                   |
| G      | likelihood-weighted pool over {leaf, A, B, C, D, E, F}                         |
| H      | likelihood-weighted pool over {leaf, A, B, F} (G minus the gaussianize chains) |
| I      | the pool of G with kernel-smoothed gaussianize                                 |

Configs C and D are the conformal production rule written as a chain: transform, apply the fence-post empirical map, stop, and read the predictive off the inverse. Config E continues through the operator. The pools use the same likelihood-weighted online ensembling as the companion paper’s forecaster, with a complexity penalty by chain depth.

A second campaign, run for this revision, adds the decomposition ladder and robustness grids reported below. Config D0 is config D with its fitted leaf replaced by the fixed  $N(0, 1)$  law, the literal conformal terminal forecast, and a full-grid variant retains every knot. Alternative fixed links replace gaussianize in the E position. Guard and resolution grids, an iid null suite, and a CRPS-weighted pool complete the set. The refetch recovers 572 of the 701 series, mostly dropping short monthly ones, and the campaign runs use those, with percentile bootstrap intervals resampling series. The dropped series carried a larger share of the operator’s log-score tax on the original run, which is part of why the tax reads 0.079 nats here against 0.108 there, the rest being data vintage. The corpus can always grow. Figure 1 summarizes the comparisons behind the paper’s three findings.

## 4.2 Stopping and inserting are different questions

In conformal practice the empirical map is terminal. One fits a model, ranks the residuals, and quotes the ranked law as the predictive. The fence-post step is where inference ends, and the design assumes the residuals arrive exchangeable, with nothing left in them worth modeling. Config D is that pattern written as a chain up to one caveat: its terminal leaf is a fitted unconditional Gaussian. Config D0 removes the caveat by quoting the fixed  $N(0,1)$  law at the operator’s output, which is the literal conformal terminal forecast. Config E reads the same operator differently: its output is a coordinate system, and whatever structure the model missed is easier to fit in those coordinates. The two readings are testable against each other. For the exact sequential ranks, exchangeable input leaves no serial rank signal for a refit to find, by Proposition 8. The interpolated coordinates carry spacing information besides, and an estimated AR(1) brings estimation noise of its own, so the comparison is informative rather than clean.

The ladder, on the 572-series campaign corpus, separates three prices. Knot compression is free: the full-grid variant differs from 39-knot D0 by 0.003 nats, compressed ahead. Refitting the terminal leaf hurts: D scores 0.027 nats below D0 and wins 3 percent of series, because the operator has already standardized its output and a fitted variance adds estimation noise. Serial refitting beats the literal stop on average: E improves on D0 by 0.019 nats (interval 0.011 to 0.029) and cuts the mean per-series CRPS ratio by 3.9 percent, winning 66 percent of series on CRPS. The mean CRPS gain is tail-driven, with a median cut of 0.3 percent, and the log-score wins are a minority. E beats D0 on a third of series, loses slightly on the rest, and gains a lot where it gains. Against the fitted-leaf D the margin is 0.046 nats with 73 percent wins.

Under the log score the oracle version of this margin is the predictive mutual information remaining in the transformed stream, the quantity the companion paper’s ledger prices (Cotton, 2026b). The measured margin is the part of it an AR(1) captures, net of estimation cost. Serial structure survives the transform, and exploiting it is a skewed trade rather than a uniform one: the conformal stopping point is defensible series by series and beaten on the corpus.

A synthetic experiment shows the mechanism in isolation. Let  $z_t = 0.8z_{t-1} + \varepsilon_t$  and observe  $y_t = \text{sign}(z_t)|z_t|^{1.5}$ , a monotone warp of an AR(1), which obscures the linear representation without removing the dependence. Stopping at gaussianize scores  $-1.56$  nats and  $0.62$  CRPS. Refitting AR(1) through it scores  $-1.07$  and  $0.41$ , beating both the stop and the same AR chain without the operator ( $-1.16$ ,  $0.42$ ). When the downstream learner is well specified only in the empirical map’s coordinates, passing through the operator beats both stopping and applying the same learner in the original coordinates.

Stopping is one question, and inserting is another. E versus D0 asks whether the rank transform is a good architectural stopping rule, and the answer is: on average, no, and on most individual series, yes. The operator’s incremental value is E versus F, the same refit without the operator, and that answer is score-dependent. Under CRPS the operator helps: E beats F on 64 percent of series and holds the best CRPS in the study. Under the log score the operator costs 0.079 nats against F on the campaign corpus and wins 9 percent of series. We return to this in Section 4.4, because the natural suspicion, that the piecewise-linear density is to blame, turns out to be wrong.

The alternative links sharpen the point. A fixed Cauchy link in the same position, the probit of the standard Cauchy distribution function, costs 0.007 nats against F on the mean, with an interval spanning zero, though it loses to F on seventy percent of series by small margins, and it captures about half of the operator’s mean CRPS gain. A sinh-arcsinh link with  $\delta = 0.5$  pays no

material tax either and wins the per-series CRPS comparison against the operator on 61 percent of series, though not on average. The rank map keeps the best mean CRPS, beating the Cauchy link on 72 percent of series, and pays for it in nats. Much of the operator’s architectural value is heavy-tailed monotone reshaping that a fixed smooth link supplies without a tax. What the fixed link cannot supply is the certificate.

### 4.3 Nesting limits the cost of admission

The raw conformal pattern C is a reasonable forecaster on this corpus, better than the location baseline on average under both scores. It is also, on a minority of series, a bad choice. Config G, the likelihood-weighted pool over all seven chains, gaussianize chains included, beats the fixed pattern C on 95 percent of series by log score, by 0.249 nats on average. On the twenty series where C is worst, with relative CRPS 1.020, the pool sits at 0.843 on the same series. The pool walks away from the production rule where it underperforms.

Admitting the pattern is approximately free on average. The average hides the shape. Config H, the same pool with the gaussianize chains removed, differs from config G by 0.012 nats and 0.4 percent of baseline CRPS, with the nesting pool G ahead on both means. Per series the picture is lopsided: G wins under a third of series by log score and just over a third by CRPS, its median difference from H is zero to four decimals, its fifth-percentile loss is under a thousandth of a nat, and its ninety-fifth-percentile gain is twenty times larger. Frequent negligible losses buy occasional real gains. Under the log score this shape is not luck. A standard Bayes-mixture regret inequality applies, and its interest here is what the subset represents: the conformal chains as a candidate block whose admission has a price.

**Proposition 9** (Nesting is cheap under the log score). *Let the pool predictive be the Bayes mixture over chains with prior  $\pi$ , and let  $q_t$  be the predictive of the renormalized sub-mixture over any subset  $Q$  of chains, with prior mass  $\pi_Q$ . Then for every realization  $y_1, \dots, y_T$ ,*

$$-\sum_{t=1}^T \log p_t(y_t) \leq -\sum_{t=1}^T \log q_t(y_t) + \log \frac{1}{\pi_Q}.$$

*Proof.* For a Bayes mixture the cumulative predictive likelihood telescopes to the prior-weighted sum of the chains’ likelihoods,  $\prod_t p_t(y_t) = \sum_c \pi_c \prod_t p_{t,c}(y_t)$ . Dropping the chains outside  $Q$  and renormalizing,  $\prod_t p_t(y_t) \geq \pi_Q \sum_{c \in Q} (\pi_c / \pi_Q) \prod_t p_{t,c}(y_t) = \pi_Q \prod_t q_t(y_t)$ . Take logarithms.  $\square$

Taking  $Q$  to be the sub-pool without the gaussianize chains, the total price of admitting them is at most  $\log(1/\pi_Q)$  nats, the prior mass of the retained block rather than of the admitted one, however the data falls, and the per-observation price vanishes. The proposition applies to the untempered Bayes pool with the depth prior applied once. The implemented pool tempers the likelihood update at rate 0.8 and applies its depth penalty per step, for adaptivity under regime change, so the measured G-versus-H comparison is empirical and does not inherit the exact bound.

Rerunning both pools as genuine Bayes mixtures, unclamped, unpruned, with the reference blend applied inside the weight updates and the depth prior applied once, reproduces the comparison on a fresh data vintage: G ahead of H by 0.012 nats on average over 572 series, the same lopsided shape, wins on one series in ten, median difference zero. The telescoping identity applies verbatim to these scores. The sub-pool prior mass is  $\pi_Q = 0.573$ , so the proposition’s worst case is 0.56 nats in total however long the series run, and the realized comparison sits

on the profitable side of it. Nothing in the reported comparison hinges on the tempering. The measured 0.012 nats in G’s favor says this corpus more than repays the admission.

No such bound is available to us under CRPS, and there the claim stays empirical. The weights update by likelihood, so the log score does the choosing and CRPS only evaluates the chosen blend. The matching experiment is a CRPS-weighted exponential pool over the same seven chains with the same prior. It beats the exact Bayes pool under the log score, by 0.021 nats per observation with wins on 76 percent of series, and the CRPS comparison is mixed: it wins 64 percent of series with median ratio 0.99, while its mean per-series CRPS ratio is 1.003, worse, because a few series blow out. The admission rule is itself a choice within the grammar, likelihood weighting is not obviously the right one, and why the log-score-driven chooser loses in its own metric is an open question we flag rather than resolve.

This is the practical content of the grammar view. A fixed conformal method is a commitment, and the companion paper prices what the commitment can cost. A grammar that contains the method as one rule converts the commitment into a weighted vote, and the vote is revised online by the scoring rule. Nesting does not avert disasters on this corpus, because the guarded operator does not produce disasters. An exact Bayes pool trims the bad tail at a bounded cumulative price, vanishing per observation. The tempered pool used here exhibits the same behavior empirically.

#### 4.4 The tested smoothing does not remove the tax

The log-score cost of the operator invites an implementation explanation: the piecewise-linear empirical map has a piecewise-constant density, and the log score punishes kinks. The kernel-smoothed  $C^1$  variant tests this directly, replacing the interpolated empirical distribution function with a smooth kernel estimate while changing nothing else. The smooth variant tracks the linear one within 0.012 nats and 0.2 percent of baseline CRPS in every position, and against config F it costs 0.113 nats where the linear variant cost 0.108, both on the original 701-series run, whose tax exceeds the campaign corpus’s 0.079 for the composition and vintage reasons of Section 4.1. The tested smoothing does not explain or remove the tax. Alternative bandwidths, adaptive tails, and splines tuned under log loss remain untested, and the remaining evidence points at resolution and tail specification.

The tax is not the kinks, and it is the same structure the companion paper measured for terminal empirical pooling. An empirical map’s resolution is bounded by its ranks, and its tails are an assumption, and the log score is exactly the score that pays for resolution and tails. Under CRPS, which pays for calibrated quantiles, the operator earns its place in the chain. Under the log score a fitted parametric leaf on the same stream is worth more. The value of the fence-post operator, like the value of residual pooling in the companion paper, depends jointly on the coordinates, the estimator, and the scoring rule. The grammar does not remove this dependence. It locates it: the scoring rule chooses among chains, and the operator is a candidate, not a default.

The probit is likewise a choice, not the point. Composing the operator with any further fixed invertible link leaves the representation content unchanged, since mutual information is invariant under invertible maps, so the transform-pooling theorem makes the link representation-irrelevant. What remains is model form. On a 120-series holdout check, a logistic link in place of the probit shifts results by hundredths of a nat under both scores, and the shift consistently favors the logistic: the Gaussian leaf pushed through the heavier-tailed link implies a heavier-tailed predictive, which fits real standardized residuals slightly better. The theorem fixes only

the representation content. A link can still matter to a restricted downstream learner through specification and estimation, and on this corpus it acts as a small tail-shape choice, an empirical observation rather than a consequence of the theorem.

## 4.5 Robustness

Three checks bound the sensitivity of these findings. An iid null suite, Gaussian, Student  $t_3$ , lognormal skew, and a two-component mixture, forty replicates each, measures how much serial signal the adaptive interpolation invents when there is nothing to find: the largest E minus D log-score delta across kinds is 0.0013 nats, with standard error of the same size, against a headline effect sixteen standard errors larger. A guard grid crossing slope caps  $\{10, 20, 40\}$  with clamps  $z_{\max} \in \{6, 8, 10\}$  on a hundred-series subsample moves config E by 0.0013 nats end to end, so the guard constants are not doing the work. A resolution grid on the same subsample finds the knot count flat in the point estimates, with 79, 159, and the full grid all within 0.003 nats of 39 and intervals spanning zero, though on a hundred series those intervals cannot exclude effects of the size the particle count shows, while raising the inverse from nine particles to nineteen is worth 0.015 nats and going further is not. The certificate’s resolution parameter is not the empirical bottleneck. The inverse’s resolution is, mildly.

## 5 One family

The grammar view collapses a scattered vocabulary. Platt scaling (Platt, 1999), isotonic calibration (Zadrozny and Elkan, 2002), temperature scaling (Guo et al., 2017), quantile recalibration (Kuleshov et al., 2018), conformal predictive systems (Vovk et al., 2019), and distributional conformal prediction (Chernozhukov et al., 2021) can all, in univariate probabilistic forecasting, be represented as monotone recalibration operators on a scalar predictive, score, or probability coordinate, differing in what they fit, what they guarantee, and where in the chain they sit. The forecasting literature’s calibration-sharpness tradition (Gneiting et al., 2007; Gneiting and Raftery, 2007) supplies the objective that orders them. In the grammar they are not competing philosophies. They are operators with types, and a pool over chains can hold several at once.

| Operator          | Coordinate | Map                   | Invertible | Certificate at output      |
|-------------------|------------|-----------------------|------------|----------------------------|
| Platt             | score      | fitted sigmoid        | yes        | none                       |
| Isotonic          | score      | fitted monotone step  | no         | none                       |
| Temperature       | logit      | fitted scale          | yes        | none                       |
| Quantile recal.   | PIT        | fitted empirical CDF  | a.e.       | asymptotic                 |
| CPS               | score      | rank map              | at knots   | finite-sample, at position |
| DCP               | PIT        | rank map              | at knots   | finite-sample, at position |
| Fence-post (here) | residual   | interpolated rank map | yes        | finite-sample, at position |

Two wrapper lines make the composition explicit. Conformal calibrators (Vovk et al., 2020) apply the rank map to the output of an arbitrary predictive system, and modular conformal calibration (Marx et al., 2022) unifies the recalibration family the table lists. Both are one level of composition, and in both the map is terminal by construction.

The time-series conformal systems nest as production rules too. Sequential predictive conformal inference (Xu and Xie, 2023) replaces the unconditional rank rung with a fitted conditional quantile model on the residual stream: the same slot, a different operator, the certificate traded for sharpness. Ensemble batch prediction intervals (Xu and Xie, 2021) keep the empirical rung

and give it a sliding-window clock. Adaptive conformal inference (Gibbs and Candès, 2021) composes an operator after the rank map, a feedback controller that moves the quoted level in response to coverage errors, and conformal PID control (Angelopoulos et al., 2023) adds a scorecaster, a serial forecaster fitted on the score stream that feeds the controller. Flow-based conformal regression (Colombo, 2024) fits an invertible reshaping upstream and applies the rank map to its output, our location and scale rungs with a richer transform.

Writing them this way shows what the grammar adds. Each system holds the rank map at a fixed position and reads out a set, an interval, or a distribution it then quotes. The one operator that literature composes after the map is a controller on the quoted level, driven by coverage. None fits a model on the map’s output, none carries predictive densities back through its inverse, and none scores the map as one operator among peers under a proper scoring rule. Those are the moves this paper adds, and the nesting is the point: every system above is a chain the pools of Section 4 could hold.

A probabilistic lifting of the same geometry appears in Snell and Griffiths (2025). If  $T_{(i)} = F(Y_{(i)})$  are the latent probability levels of the ordered observations, the spacings of the  $T_{(i)}$  are  $\text{Dirichlet}(1, \dots, 1)$ , and the plotting positions  $i/(n+1)$  used here are their means. The present operator fixes the mean positions and obtains a deterministic coordinate transform. Snell and Griffiths retain the full spacing distribution and push a loss integral through it, and standard conformal risk control is recovered by replacing the random spacings with their expectations. In grammar terms theirs is a measure-valued operator on the loss coordinate, emitting a distribution over risk bounds rather than a transformed observation, and it is complementary to the observation-coordinate transforms studied here. The two demotions agree: conformal prediction is what remains of an empirical rank geometry after one functional of it is retained and the rest discarded.<sup>2</sup>

Conformal prediction keeps a distinguished place in this family. The fence-post operator is the member whose output carries a finite-sample, distribution-free guarantee at its position in the chain (Proposition 2). No fitted recalibration offers that. The grammar view does not diminish the guarantee. It names what the guarantee attaches to, and it frees the forecaster to keep transforming afterward.

## 6 Conclusion

Conformal prediction is one production rule in a grammar of composable transforms: model, residual, fence-post empirical map, inverse. Interpolated, the fence-post map is an ordinary invertible online transform, and the rule’s distinguishing stop, quoting the empirical law where it stands, becomes a choice the scoring rule can overrule. On economic series, a serial model added after the operator outscores the literal terminal stop on average, a pool that admits the pattern trims its worst cases, at a price bounded for the exact Bayes pool and negligible empirically for the tempered one, and the operator’s remaining cost under the log score is unexplained by the tested smoothing and mostly matched by a fixed heavy-tailed link, which pays no tax and carries no certificate.

The operator is exceptional as a certificate and ordinary as an architecture, and the certificate

---

<sup>2</sup>One typing caution. The Dirichlet law is a sampling statement about the latent spacings, and reading it as a data-conditional posterior requires an argument about the prior over  $F$  that the sampling statement does not supply. Sampling-law, posterior-law, and prequential certificates are different types, and the table should keep them apart.

it carries is priced in its computational budget: retained-knot resolution, tie mass, and guard displacement set the calibration error. The typing question is the theory this view opens. An operator’s guarantee attaches to a position in a chain. Which properties survive composition is a discipline we have only begun here, with one operator, a few propositions, and the first rows of a table, and the rest of the table is open.

## A Reproducibility

Every number in Section 4 is produced by one of two runners in the `benchmarks` directory of the `skaters` repository. The original run is `gaussianize_chain.py`, with per-series results frozen in `gaussianize_chain.csv`, and it also produces the synthetic warp numbers. The campaign runner is described at the end of this appendix. The configuration is as follows. The location operator is an EMA transform with rate 0.1. The scale operator is an EWMA standardization with rate 0.05. The AR(1) operator estimates its coefficient by recursive least squares with forgetting factor 0.99 and ridge 1.0. The gaussianize operator uses at most 39 knots at equal mass from the full sorted history of its inputs (exact order statistics before the slope guard, which can displace knots when it binds; the buffer is unbounded), a warmup of 30 observations with expanding-moment standardization before it, and a 9-particle equal-probability approximation for the inverse. The kernel-smoothed variant uses a Gaussian-kernel distribution function at Silverman bandwidth, represented by a monotone  $C^1$  cubic with Fritsch–Butland tangents and inverted by Newton’s method. The guards are the inverse-slope cap at twenty times the interquartile inverse slope and the downstream-update clamp  $|z| \leq 8$ .

The pools are tempered likelihood mixtures over the candidate chains of the design table, with tempering rate 0.8, a complexity penalty of 0.005 per unit of chain depth applied to the online weights, and at most 20 mixture components. Scoring starts after a burn-in of 300 observations. The log score mixes each predictive with weight 0.01 toward an expanding-window Gaussian reference and is otherwise exact. CRPS is computed in closed form from the Gaussian-mixture predictive and normalized per series by config A. The corpus is 701 FRED change series with at least 600 observations and repeat fraction below five percent, listed in the frozen results file. Series were downloaded from FRED on 2 August 2026 with observation start 2005-01-01, and non-numeric observations are dropped at parse. The runner does not log guard activations separately. The one recorded activation episode is the warmup event described in Section 2.

The decomposition ladder, alternative links, null suite, guard and resolution grids, and CRPS pool are produced by `benchmarks/grammar_campaign.py` in the same repository, with per-series results frozen in the `grammar_campaign_*.csv` files. The campaign corpus is the 572 of the 701 series whose full history could be refetched on 3 August 2026. Config D0 ends in a fixed  $N(0,1)$  leaf. The fixed links use the same nine-particle inverse pushback as gaussianize: Yeo–Johnson at  $\lambda \in \{0.5, 1.5\}$ , sinh-arcsinh at  $\delta \in \{0.5, 2\}$ , and the Cauchy link  $\Phi^{-1}(F_{\text{Cauchy}})$ . The null suite draws series of length 1200 and forty replicates per kind. The CRPS pool uses exponential weights with rate 1 on cumulative CRPS and the exact-Bayes prior. The guard and resolution subsamples take every fifth series alphabetically, truncated to one hundred, a deterministic rule that never samples the last seventy-odd series. The deletion-guard comparison is the `delguard` subcommand, using the operator’s `guard` option, on the full campaign corpus. The exact Bayes pools are produced by `grammar_exact_bayes.py`, which updates weights by exact blended likelihoods without clamping or pruning and reports CRPS from a pruned copy of the predictive, the log score being exact. Bootstrap intervals are percentile intervals resampling

series, 2000 draws, seed 0.

## References

- Anastasios N. Angelopoulos, Emmanuel J. Candès, and Ryan J. Tibshirani. Conformal PID control for time series prediction. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- Andreas Auer, Martin Gauch, Daniel Klotz, and Sepp Hochreiter. Conformal prediction for time series with modern hopfield networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2303.12783.
- Henrik Boström, Ulf Johansson, and Tuwe Löfström. Mondrian conformal predictive distributions. In *Conformal and Probabilistic Prediction and Applications (COPA)*, PMLR, volume 152, pages 24–38, 2021.
- George E. P. Box and David R. Cox. An analysis of transformations. *Journal of the Royal Statistical Society, Series B*, 26(2):211–243, 1964.
- Victor Chernozhukov, Kaspar Wüthrich, and Yinchu Zhu. Distributional conformal prediction. *Proceedings of the National Academy of Sciences*, 118(48):e2107794118, 2021.
- Nicolo Colombo. Normalizing flows for conformal regression. In *Proceedings of the 40th Conference on Uncertainty in Artificial Intelligence*, volume 244 of *Proceedings of Machine Learning Research*, 2024.
- Peter Cotton. *Microprediction: Building an Open AI Network*. MIT Press, 2022.
- Peter Cotton. An algebra of prediction-rewarding mechanisms. <https://mechanisms.microprediction.org/papers/composition-and-the-algebra-of-mechanisms.html>, 2026a.
- Peter Cotton. An empirical study of the conformal information gap. SSRN preprint, 2026b.
- Peter Cotton. Multi-stage solicitation of probability distributions. <https://mechanisms.microprediction.org/papers/multi-stage-solicitation.html>, 2026c.
- Peter Cotton. Transforms all the way down: Automatic online distributional forecasting by conjugation. <https://github.com/microprediction/skaters/blob/main/papers/skaters-jss.pdf>, 2026d. Software: <https://github.com/microprediction/skaters>.
- Peter Cotton. skaters: Online distributional time-series forecasting by conjugation, in Python and JavaScript. <https://github.com/microprediction/skaters/blob/main/papers/joss/paper.md>, 2026e. Software paper, under review at the Journal of Open Source Software.
- David R. Cox and E. Joyce Snell. A general definition of residuals. *Journal of the Royal Statistical Society, Series B*, 30(2):248–265, 1968.
- Florence N. David and Norman L. Johnson. The probability integral transformation when parameters are estimated from the sample. *Biometrika*, 35(1/2):182–190, 1948.

- A. Philip Dawid. Present position and potential developments: some personal views. Statistical theory: the prequential approach. *Journal of the Royal Statistical Society, Series A*, 147(2): 278–292, 1984.
- Francis X. Diebold, Todd A. Gunther, and Anthony S. Tay. Evaluating density forecasts with applications to financial risk management. *International Economic Review*, 39(4):863–883, 1998.
- Rina Foygel Barber, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. The limits of distribution-free conditional predictive inference. *Information and Inference*, 10(2):455–482, 2021.
- Francis Galton. Statistics by intercomparison, with remarks on the law of frequency of error. *Philosophical Magazine, Fourth Series*, 49:33–46, 1875.
- Isaac Gibbs and Emmanuel J. Candès. Adaptive conformal inference under distribution shift. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 34, pages 1660–1672, 2021.
- Tilmann Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007.
- Tilmann Gneiting, Fadoua Balabdaoui, and Adrian E. Raftery. Probabilistic forecasts, calibration and sharpness. *Journal of the Royal Statistical Society: Series B*, 69(2):243–268, 2007.
- Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, pages 1321–1330, 2017.
- Allen Hazen. Storage to be provided in impounding reservoirs for municipal water supply. *Transactions of the American Society of Civil Engineers*, 77:1539–1640, 1914.
- Volodymyr Kuleshov, Nathan Fenner, and Stefano Ermon. Accurate uncertainties for deep learning using calibrated regression. In *Proceedings of the 35th International Conference on Machine Learning*, pages 2796–2804, 2018.
- Jing Lei and Larry Wasserman. Distribution-free prediction bands for non-parametric regression. *Journal of the Royal Statistical Society, Series B*, 76(1):71–96, 2014.
- Jing Lei, Max G’Sell, Alessandro Rinaldo, Ryan J. Tibshirani, and Larry Wasserman. Distribution-free predictive inference for regression. *Journal of the American Statistical Association*, 113(523):1094–1111, 2018.
- Charles Marx, Shengjia Zhao, Willie Neiswanger, and Stefano Ermon. Modular conformal calibration. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, 2022.
- Hans A. Panofsky and Glenn W. Brier. *Some Applications of Statistics to Meteorology*. Pennsylvania State University, 1968.

- George Papamakarios, Eric Nalisnick, Danilo Jimenez Rezende, Shakir Mohamed, and Balaji Lakshminarayanan. Normalizing flows for probabilistic modeling and inference. *Journal of Machine Learning Research*, 22(57):1–64, 2021.
- Egon S. Pearson. The probability integral transformation for testing goodness of fit and combining independent tests of significance. *Biometrika*, 30(1/2):134–148, 1938.
- John C. Platt. Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. In *Advances in Large Margin Classifiers*, pages 61–74. MIT Press, 1999.
- Alfréd Rényi. Théorie des éléments saillants d’une suite d’observations. *Annales scientifiques de l’Université de Clermont-Ferrand 2*, 8:7–13, 1962.
- Yaniv Romano, Evan Patterson, and Emmanuel J. Candès. Conformalized quantile regression. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 32, 2019.
- Murray Rosenblatt. Remarks on a multivariate transformation. *The Annals of Mathematical Statistics*, 23(3):470–472, 1952. doi: 10.1214/aoms/1177729394.
- Jake C. Snell and Thomas L. Griffiths. Conformal prediction as Bayesian quadrature. arXiv:2502.13228, 2025.
- Esteban G. Tabak and Eric Vanden-Eijnden. Density estimation by dual ascent of the log-likelihood. *Communications in Mathematical Sciences*, 8(1):217–233, 2010.
- Bartel L. van der Waerden. Order tests for the two-sample problem and their power. *Indagationes Mathematicae*, 14:453–458, 1952.
- Vladimir Vovk. Conditional validity of inductive conformal predictors. In *Proceedings of the Asian Conference on Machine Learning*, volume 25 of *Proceedings of Machine Learning Research*, pages 475–490, 2012.
- Vladimir Vovk, Ilia Nouretdinov, and Alex Gammerman. Testing exchangeability on-line. In *Proceedings of the Twentieth International Conference on Machine Learning*, pages 768–775, 2003.
- Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, New York, 2005.
- Vladimir Vovk, Jieli Shen, Valery Manokhin, and Min-ge Xie. Nonparametric predictive distributions based on conformal prediction. *Machine Learning*, 108(3):445–474, 2019.
- Vladimir Vovk, Ivan Petej, Paolo Toccaceli, Alexander Gammerman, Ernst Ahlberg, and Lars Carlsson. Conformal calibrators. In *Proceedings of the Ninth Symposium on Conformal and Probabilistic Prediction and Applications*, volume 128 of *Proceedings of Machine Learning Research*, pages 84–99, 2020.
- Chen Xu and Yao Xie. Conformal prediction interval for dynamic time-series. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 11559–11569, 2021.

- Chen Xu and Yao Xie. Sequential predictive conformal inference for time series. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 38707–38727, 2023.
- In-Kwon Yeo and Richard A. Johnson. A new family of power transformations to improve normality or symmetry. *Biometrika*, 87(4):954–959, 2000.
- Bianca Zadrozny and Charles Elkan. Transforming classifier scores into accurate multiclass probability estimates. In *Proceedings of the Eighth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 694–699, 2002. doi: 10.1145/775047.775151.