

Moving the Slots, Not the Balls: Conformal Prediction as the Contragredient of Herding

Peter Cotton
peter.cotton@microprediction.com

July 2026

Abstract

Split conformal prediction is herding run on the other side of the pairing. Herding moves particles to balance fixed tests, earning $1/n$ from a bounded weight ledger; conformal holds the scores fixed and transports the test label across slots, where the rank count is already balanced — exactly $\lceil(1 - \alpha)(n + 1)\rceil$ of the $n + 1$ placements are covered, for any scores. The empirical quantile is the matching feedback controller, holding the accumulated acceptance-rate error below one unit, and the switch-coefficient bound for dependent series (Barber and Pananjady, 2026) is the same ledger: τ/n for deleted slots plus a total-variation charge for failed rotation invariance. Nothing concentrates, which is why both rates beat $\sqrt{n}$. Exchangeability still does all the probabilistic work.

1 Two mechanisms

Herding moves the balls. Fix a feature map $\phi : \mathcal{X} \rightarrow \mathcal{H}$ and a target expectation $\mu = \mathbb{E}_P \phi(X)$. Herding (Welling, 2009) generates points deterministically,

$$x_t = \arg \max_{x \in \mathcal{X}} \langle w_{t-1}, \phi(x) \rangle, \quad w_t = w_{t-1} + \mu - \phi(x_t), \quad (1)$$

so that the weight vector $w_t = w_0 + t\mu - \sum_{s \leq t} \phi(x_s)$ is, up to initialization, t times the running discrepancy between the target mean and the empirical feature mean. If the dynamics keep $\|w_t\|$ bounded — guaranteed in finite dimension under mild conditions, and more delicately in a general RKHS (Chen et al., 2010; Bach et al., 2012) — then

$$\left\| \frac{1}{n} \sum_{t=1}^n \phi(x_t) - \mu \right\| = \frac{\|w_n - w_0\|}{n} = O\left(\frac{1}{n}\right), \quad (2)$$

a quasi-Monte Carlo rate, against the $O(n^{-1/2})$ of i.i.d. sampling. The tests ϕ are held fixed; the *state-space locations* of the particles are what the algorithm manipulates. Bach et al. (2012) identify (1) with the Frank-Wolfe algorithm on the marginal polytope, which is where the boundedness conditions live; Chen et al. (2010) read it as greedy quasi-Monte Carlo.

Conformal moves the slots. Split conformal prediction (Vovk et al., 2005; Lei et al., 2018) computes conformity scores $S_1, \dots, S_n$ on a calibration set, a score S_{n+1} on the test point, and declares coverage when S_{n+1} falls below the inflated empirical quantile $Q_{1-\alpha} =$ the k -th smallest of $S_1, \dots, S_n$, with $k = \lceil(1 - \alpha)(n + 1)\rceil$. Nothing in the procedure moves a score. The proof of validity instead moves a *label*. Write $S = (S_1, \dots, S_{n+1}) \in \mathbb{R}^{n+1}$ for the realized scores — the balls, resting where

they fell — and let e_i^* denote the coordinate functional that reads slot i . The coverage event is a statement about the pairing $\langle e_{n+1}^*, S \rangle$ and the symmetric function $Q_{1-\alpha}$ of the remaining coordinates. The proof asks: what if the designation “test” had pointed at slot i instead?

Observation 1 (The rank count is already balanced). *Let $s_1, \dots, s_{n+1}$ be distinct reals and $k = \lceil (1-\alpha)(n+1) \rceil$. For each placement $i \in \{1, \dots, n+1\}$, say the placement is covered when s_i is at most the k -th smallest of the other n values. Then the number of covered placements is exactly k : placement i is covered if and only if the rank of s_i among all $n+1$ values is at most k . In particular the covered fraction is $k/(n+1) \geq 1-\alpha$, deterministically, for any score vector whatsoever.*

This is the whole finite-sample content of the guarantee. No mass is transported, no point is steered; the balance that herding must manufacture by moving particles is already present in the rank structure, and the proof merely rotates the pointer to exhibit it. Probability enters once, at the end: exchangeability says the realized designation (slot $n+1$) is indistinguishable from a uniform draw over the $n+1$ placements, so the deterministic fraction becomes a probability.

2 The contragredient

The duality can be said precisely. The symmetric group acts on score vectors through permutation matrices, $\rho(\pi)S = P_\pi S$; the contragredient (dual) representation acts on covectors, $\rho^*(\pi)f = f \circ \rho(\pi^{-1})$, so that the pairing is invariant:

$$\langle \rho^*(\pi) e_{n+1}^*, \rho(\pi) S \rangle = \langle e_{n+1}^*, S \rangle.$$

Every exchangeability proof of conformal validity can be run on either side of this pairing: *either* rearrange the data, $S \mapsto P_\pi S$, holding the test functional fixed, *or* hold the data fixed and pull the functional back, $e_{n+1}^* \mapsto \rho^*(\pi^{-1})e_{n+1}^* = e_{\pi(n+1)}^*$. Because the quantile threshold is a symmetric — that is, invariant — function of the score multiset, the two descriptions agree event by event. The textbook proof is usually narrated in the first voice (“swap the test point with a calibration point”); Observation 1 is the second voice, and the second voice is the one in which the mechanism is visible: the orbit being averaged is the orbit of the *functional*, $\{e_i^*\}_{i=1}^{n+1}$, not an orbit of datasets.

One honesty clause first: permutation representations are orthogonal, so the contragredient is *equivalent* to the original representation; the term is apt not because the dual carries different structure but because it names which side of the pairing the action sits on — the covector side.

Where, then, is the steering? An earlier draft of this note claimed conformal performs none; that is wrong, and the correction sharpens the duality. Write $p = 1 - \alpha$ and $h_q(s) = \mathbf{1}\{s \leq q\}$ for the acceptance test at threshold q . Choosing the calibration quantile $q_n = S_{(k)}$ with $k = \lceil p(n+1) \rceil$ enforces, up to ties, the identity

$$0 \leq \sum_{i=1}^n (h_{q_n}(S_i) - p) < 1 \quad (\text{up to the } \pm 1 \text{ bookkeeping of the } (n+1)\text{-inflation}), \quad (3)$$

and (3) is unmistakably a feedback-control identity: the realized scores determine how the cutoff moves, and the cutoff moves so as to keep the accumulated acceptance-rate error bounded by one unit — the exact analogue of herding’s bounded weight $\|w_n\|$, and the same division by n yields the same $1/n$. The difference between the two mechanisms is therefore not steering versus no steering but *which member of the pairing* $\langle h, \nu \rangle = \int h d\nu$ is controlled: herding moves the empirical measure $\nu_n = \frac{1}{n} \sum_i \delta_{x_i}$ against fixed tests $h = \phi$; conformal calibration moves the test h_{q_n} against the fixed empirical measure of scores. There are, in fact, two slot movements in a conformal argument, and it is worth keeping them separate:

steering: the acceptance boundary q_n moves until the empirical rank count balances;
transport: the proof moves the designation “test” among the slots (Observation 1).

Under exchangeability the transport is exact — the pointer is equally entitled to indicate any slot — and the guarantee is the steering identity carried to the future slot intact. The residual asymmetry with herding is only this: conformal does not control the generation of the next observation. The steering is real, but it lives in the dual space of thresholds, acceptance sets, and index labels rather than in the primal space of generated points:

herding moves the observations to balance the measure;
conformal moves the acceptance boundary and the observation *labels*
to balance — and then transport — an empirical rank count.

3 The $1/n$ ledger, and dependence as damaged slots

Both mechanisms produce $1/n$ rates, and for the same bookkeeping reason: an error that accumulates *boundedly* is divided by the sample size.

$$\begin{aligned} \text{herding:} \quad \text{error} &\leq \frac{\text{bounded cumulative feature discrepancy } \|w_n\|}{n}, \\ \text{conformal:} \quad \text{coverage damage} &\leq \frac{\text{number of damaged or deleted ranks}}{n}. \end{aligned}$$

The second line is not a metaphor; it is the proof technique of Barber and Pananjady (2026) for split conformal on a dependent series. There, exact slot symmetry is unavailable: rotating the test designation to slot k is no longer distribution-preserving, because neighbouring scores remember one another. The repair is to delete a buffer of τ slots around the splice before rotating. Deleting τ of n entries moves the k -th order statistic by at most τ ranks — order statistics are insertion/deletion stable — so the quantile level shifts by at most τ/n ; and the residual failure of rotation invariance is charged at its total-variation price, the switch coefficient $\bar{\Psi}_\tau$, bounded by $2\beta(\tau)$ for a stationary β -mixing sequence. The coverage floor

$$\mathbb{P}(\text{cover}) \geq 1 - \alpha - \min_{\tau} \left\{ \underbrace{\frac{\tau}{n+1}}_{\text{slots sacrificed}} + \underbrace{\bar{\Psi}_\tau}_{\text{failed invariance}} \right\}$$

is index-space low-discrepancy accounting — in the decomposition above, the steering term is the exact empirical rank correction and the switch term is the price of transporting that correction to the future slot — and nothing in it concentrates: no variance appears, no square root can. The contrast is instructive — when the same problem is attacked by blocking and empirical-process concentration (Oliveira et al., 2024), the rate degrades to the Monte Carlo $\sqrt{\tau/n}$, precisely because the counting structure has been traded away for a maximal inequality. Herding’s advantage over i.i.d. sampling and the switch coefficient’s advantage over blocking are the same advantage.

The quasi-Monte Carlo reading closes the loop. Chen et al. (2010) present herding as a greedy low-discrepancy construction: choose the next point to keep the running discrepancy bounded. The rank sequence of any sample is, tautologically, a perfectly low-discrepancy sequence over its own index set — every rank appears exactly once — and Observation 1 is the Koksma–Hlawka-style consequence: a test functional averaged over slots inherits the balance of the ranks. Conformal prediction is what you get when the low-discrepancy sequence comes for free and only the pointer needs transporting. The quadrature reading of conformal prediction is not ours alone: Snell and

Griffiths (2025) recast the coverage computation as Bayesian quadrature over the calibration scores, recovering the frequentist guarantee as a special case of a posterior over test loss. The slot-transport view is the deterministic, index-space face of the same coin: quadrature weights over scores on one side, a rotating evaluation functional over slots on the other.

4 What the duality does and does not explain

It explains the finite-sample ledger: why the guarantee is exact at finite n , why corrections enter linearly in the number of damaged slots, and why dependence is a tax of size (dependence range)/(calibration size) rather than its square root. It does not explain, and we do not claim it explains, the statistical content: the step from “ k of $n + 1$ placements are covered” to “the realized placement is covered with probability $k/(n + 1)$ ” is exchangeability, or its quantitative surrogate under mixing, and no amount of slot bookkeeping supplies it. Nor does dual-space steering create information: moving thresholds and labels over fixed scores can re-level coverage exactly, but the conditional sharpness of the resulting intervals is fixed by the scores themselves — the information-gap point made elsewhere on this site (Cotton). A natural question is whether the steering can be aimed: a herding-like feedback in index space that retains or reweights calibration slots to balance the rank count *conditionally*, not just marginally. Localized and weighted conformal methods are exactly such aimed controllers, and the same information-theoretic ceiling applies to them. We leave the speculation at one sentence.

References

- F. Bach, S. Lacoste-Julien, and G. Obozinski. On the equivalence between herding and conditional gradient algorithms. In *Proceedings of the 29th International Conference on Machine Learning (ICML)*, pages 1355–1362, 2012.
- R. F. Barber and A. Pananjady. Predictive inference for time series: why is split conformal effective despite temporal dependence?, 2026. Preprint, arXiv:2510.02471.
- Y. Chen, M. Welling, and A. Smola. Super-samples from kernel herding. In *Proceedings of the 26th Conference on Uncertainty in Artificial Intelligence (UAI)*, pages 109–116, 2010.
- P. Cotton. Marginally useful: Formalizing the information gap in conformal prediction. Companion paper.
- J. Lei, M. G’Sell, A. Rinaldo, R. J. Tibshirani, and L. Wasserman. Distribution-free predictive inference for regression. *Journal of the American Statistical Association*, 113(523):1094–1111, 2018.
- R. I. Oliveira, P. Orenstein, T. Ramos, and J. V. Romano. Split conformal prediction and non-exchangeable data. *Journal of Machine Learning Research*, 25, 2024.
- J. C. Snell and T. L. Griffiths. Conformal prediction as bayesian quadrature. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, 2025. arXiv:2502.13228.
- V. Vovk, A. Gammerman, and G. Shafer. *Algorithmic Learning in a Random World*. Springer, 2005.
- M. Welling. Herding dynamical weights to learn. In *Proceedings of the 26th International Conference on Machine Learning (ICML)*, pages 1121–1128, 2009.